TRAIL OFBITS

# Trusted Publishing: Lessons from PyPI

William Woodruff

#### Introduction

### Hello!

### William Woodruff (william@trailofbits.com)

- open source group engineering director @ trail of bits
- long-term OSS contributor (Homebrew, LLVM, Python) and maintainer (pip-audit, sigstore-python)
- @yossarian@infosec.exchange

### **Trail of Bits**

- ~150 person cybersecurity engineering and auditing consultancy
- specialities: cryptography, compilers, program analysis research, "supply chain", OSS package management, general high assurance software development

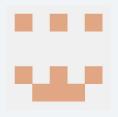




### Some thank-yous

- This work wouldn't have happened without Dustin Ingram and the GOSST team's vision for improving PyPI's security!
- The other maintainers of PyPI (Donald, Ee, Mike) all reviewed this work or otherwise made it possible
- Other members of PyPA for being early testers and ensuring a smooth rollout
  - Special thanks to Sviatoslav Sydorenko
- PyCA maintainers (Paul and Alex) for being guinea pigs and providing early feedback on usability





### Magic

# What's wood gight with is his cture? picture?



```
- name: publish
 uses: pypa/gh-action-pypi-publish@master
 with:
                                                   /v1
    user: __token__
    password: ${{ secrets.PYPI_TOKEN }}
```

#### Version 2.6.1 created

Added by: OpenID created token URL: https://github.com/pypa/pip-

audit/commit/d4242095300357730e1510ef2db837cf1b1142f1

### **Background**

- Pronounced >---
- The primary package index for the Python ecosystem
  - ~500K projects, ~5M releases, ~9M files, ~750K users (i.e. packagers)
  - ~20 **billion** downloads per month (August 2023)
- Rewritten in ~2017, security features added since
  - 2018/19: API tokens, TOTP and WebAuthn, security event logging for users and projects
  - 2019/20: Malware scanning
  - 2020/21: Vulnerability feeds, GitHub secret scanning integration
  - 2022/23: Trusted publishing (you are here!)



### **Background**

# Uploading to PyPI: then and ~now

- < 2019: username/password authentication for uploads
  - No credential separation: same user/pass could **modify all projects** under the same user!
    - As well as log into PyPI and do normal account admin things
  - No straightforward revocation: compromise means full account recovery needed!
  - No security events: attacker who steals creds can remain (relatively) stealthy!
- >= 2019: Macaroon-based API tokens
  - Configurable scopes (per-user, per-project)
    - Per-user means "all projects," not "can modify the user's profile"!
  - Integrated into security events + GitHub secret scanning
  - Still need to be manually configured + revoked on compromise
  - Too easy to over-scope (developer confusion, fatigue)
    - Chicken-and-egg problems with new projects (can't scope a token for a nonexistent project)

API tokens are a major improvement; let's do even better!

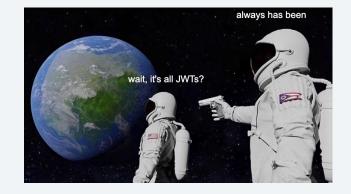
## Doing better than API tokens

API tokens are still somewhat manual; doing better means automation.

#### What we want:

- No more manual token transfers: CI/CD should receive credentials automatically
- Maximum scoping: credentials should be scoped down to the smallest unit of work they're intended to perform, and should automatically expire once that work is done
- Misuse resistance: users can't leak credentials that don't exist!





- CI/CD providers like GitHub support machine identities through OIDC
  - These credentials are strongly bound to the repository + workflow that made them
  - Can be verified by any third-party service using OIDC Discovery!
- OIDC credentials are short-lived + scoped to an intended audience
  - Fewer compromise opportunities + no domain contamination!
  - Service B won't accept credentials made for Service A
- Already widely applied to other services (GH  $\longleftrightarrow$  AWS, GCP, etc.)
  - Why not PyPI too?

# OpenID Connect for PyPI publishing

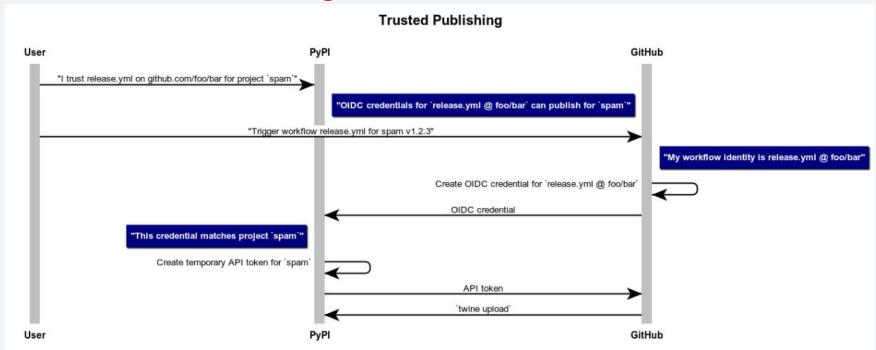
### The rough idea:

- Ahead of time: Users configure a trust relationship between a specific OIDC provider (e.g. GitHub) and their PyPI project
  - The trust relationship itself isn't secret, so no potential leaks here!
  - For GitHub: user/repo slug, workflow name (e.g. release.yml), optional environment name
  - We call this relationship the "trusted publisher"
- During publishing: OIDC provider creates an OIDC credential
  - PyPI accepts that credential, verifies it, and exchanges it for a short-lived PyPI API token scoped for the project
  - Package publishing (e.g. through twine) continues as normal, none the wiser!

OWNET The GitHub organization name or GitHub username that owns the repository Repository name (required) The name of the GitHub repository that contains the publishing workflow Workflow name (required) Workflow.yml The filename of the publishing workflow. This file should exist in the _github/workflows/ directory in the repository configured above.	Read more about GitHub Actions's OpenID Connect su  Owner (required)	
The GitHub organization name or GitHub username that owns the repository  Repository name (required)  repository  The name of the GitHub repository that contains the publishing workflow  Workflow name (required)  Workflow.yml  The filename of the publishing workflow. This file should exist in the   _github/workflows/  directory in the repository configured above.  Environment name (optional)  release  The name of the <u>GitHub Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't		
Repository name (required)  repository  The name of the GitHub repository that contains the publishing workflow  Workflow name (required)  Workflow.yml  The filename of the publishing workflow. This file should exist in the _github/workflows/ directory in the repository configured above.  Environment name (optional)  release  The name of the GitHub Actions environment that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	owner	
repository  The name of the GitHub repository that contains the publishing workflow  Workflow name (required)  Workflow.yml  The filename of the publishing workflow. This file should exist in the _github/workflows/ directory in the repository configured above.  Environment name (optional)  release  The name of the <u>GitHub Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	The GitHub organization name or GitHub username that own repository	s the
The name of the GitHub repository that contains the publishing workflow name (required)  Workflow.yml  Workflow.yml  Workflow.yml  Workflow.yml  The filename of the publishing workflow. This file should exist in the i_github/workflows/  directory in the repository configured above.  Environment name (optional)  release  The name of the GitHub Actions environment that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	Repository name (required)	
Workflow name (required)  Workflow.yml  The filename of the publishing workflow. This file should exist in the [.github/workflows/] directory in the repository configured above.  Environment name (optional)  release  The name of the <u>Github Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	repository	
workflow.yml  The filename of the publishing workflow. This file should exist in the i_github/workflows/  directory in the repository configured above.  Environment name (optional)  release  The name of the Github Actions environment that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated bublishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	The name of the GitHub repository that contains the publishi workflow	ng
The filename of the publishing workflow. This file should exist in the [.github/workflows/] directory in the repository configured above.  Environment name (optional)  release  The name of the Github Actions environment that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	Workflow name (required)	_
the _github/workflows/ directory in the repository configured above.  Environment name (optional)  release  The name of the <u>Github Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your epository has maintainers with commit access who shouldn't	workflow.yml	
release  The name of the <u>GitHub Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is strongly encouraged, especially if your repository has maintainers with commit access who shouldn't		
The name of the <u>GitHub Actions environment</u> that the above workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated publishing environment is <b>strongly</b> encouraged, <b>especially</b> if your epository has maintainers with commit access who shouldn't		
workflow uses for publishing. This should be configured under the repository's settings. While not required, a dedicated bublishing environment is strongly encouraged, especially if your repository has maintainers with commit access who shouldn't	Environment name (optional)	

Add a new publisher

# Trusted publishing: the bird's-eye view

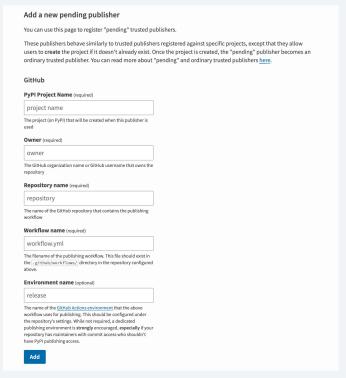


### What about nonexistent projects?

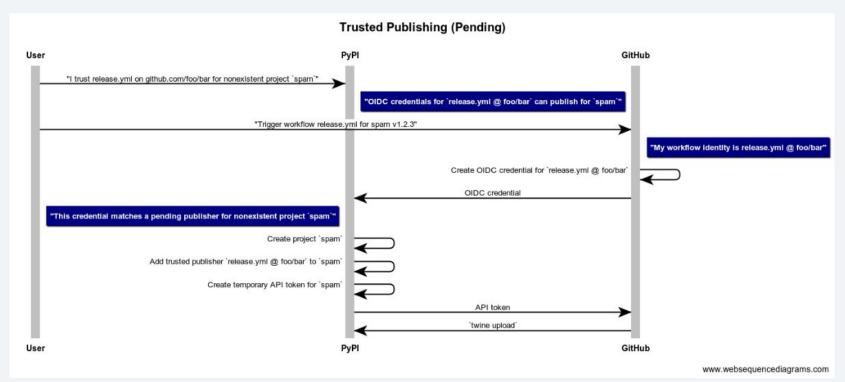
This first approach doesn't solve the "chicken-and-egg" problem with API tokens: the project still needs to exist to register a trusted publisher to it!

We solve this with "pending publishers": registered similarly to a trusted publisher, but associated with a user instead.

- Contains the name of the project that will be created
- On first use, the project is created (in an empty state) and the "pending" publisher is reified into a full "trusted publisher"



# Pending publishers: the bird's-eye view



### It works!

With this scheme, we achieve all of our security goals:

- All credentials are temporary and self-expiring
- All credentials are minimally scoped (no user scopes)
- Users only perform configuration once (initial trusted setup)
- All configuration is over public information (no private metadata)
- No more chicken-and-egg ("pending" publishers transition seamlessly to full publishers once used)

We also solve supply chain problems in the process:

- Flattening of state: the source repository itself becomes the "ground truth"
  - Maintenance transitions: projects can transition maintainers without playing "who owns the credential"



# nitty-gritty time

lessons for other possible implementations



### Lesson #1: the data model is unintuitive

"One trusted publisher per package, how hard could it be?"

### Realities:

- Multiple logical projects live under the same logical publisher
  - Example: GitHub monorepo with multiple PyPI projects
- Multiple logical publishers are responsible for a single project
  - Example: Single PyPI project with multiple arch-specific publishing workflows
    - Users ideally wouldn't do this, but we want to encourage adoption!

Conclusion: trusted publishing is actually many-many; this may have surprising complexity implications for the ecosystem you're adding it to!

```
class OIDCPublisherProjectAssociation(db.Model):
    __tablename__ = "oidc_publisher_project_association"
    oidc_publisher_id = Column(
        UUID(as_uuid=True),
        ForeignKey("oidc_publishers.id"),
        nullable=False,
        primary_key=True,
    project_id = Column(
        UUID(as_uuid=True), ForeignKey("projects.id"), nullable=False, primary_key=True
```

# Lesson #2: OIDC is very narrowly standardized

"They're just JWTs under the hood, how different could they be?"

#### Realities:

Individual providers have wide latitude in claim availability/format

- Only basic things can be assumed to be universal: iss, exp, aud, etc.
- Each new provider needs to be carefully inspected to determine which parts of the OIDC credential constitute sufficient trusted metadata
  - This requires in-depth knowledge of the provider's internals/behavior, e.g. which users are entitled to run GitHub workflows within a particular repository!

#### **Conclusions:**

- Adding new trusted publisher providers (GitLab, etc.) is time intensive; ecosystems should prioritize the providers they see used the most (GitHub for PyPI)
- Supporting multiple providers = more data model complexity!

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NO as described in RFC 2119 [RFC2119].

```
class OIDCPublisherMixin:
    """

A mixin for common functionality between all OIDC publishers, including
    "pending" publishers that don't correspond to an extant project yet.
    """

# Each hierarchy of OIDC publishers (both `OIDCPublisher` and
    # `PendingOIDCPublisher`) use a `discriminator` column for model
    # polymorphism, but the two are not mutually polymorphic at the DB level.
    discriminator = Column(String)

# A map of claim names to "check" functions, each of which
    # has the signature `check(ground-truth, signed-claim, all-signed-claims) -> bool`.
    __required_verifiable_claims__: dict[str, CheckClaimCallable[Any]] = dict()

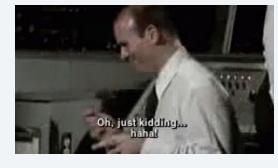
# Simlar to __verifiable_claims__: dict[str, CheckClaimCallable[Any]] = dict()
```

```
class OIDCPublisher(OIDCPublisherMixin, db.Model):
    __tablename__ = "oidc_publishers"

projects = orm.relationship(
    Project,
    secondary=OIDCPublisherProjectAssociation.__table__, # type: ignore
    backref="oidc_publishers",
)
macaroons = orm.relationship(Macaroon, cascade="all, delete-orphan", lazy=True)

__mapper_args__ = {
    "polymorphic_identity": "oidc_publishers",
    "polymorphic_on": OIDCPublisherMixin.discriminator,
}
```

### Lesson #2.5: OIDC varies wildly within IdPs



"Every trusted publisher through e.g. GitHub should look basically the same"

### Realities:

- OIDC identities vary wildly even within a provider: GitHub has special claims for reusable workflows, claims for different CI event types, etc.

  • Providers like to change their claims without telling anyone!
- Differences between these claims can't be paved over without (1) excluding some users or (2) ignoring some claims that might be important!

### **Conclusion:**

Supporting every possible configuration of a trusted publisher is hard + trying to do so opens up a lot of potential logic errors!

```
@staticmethod
def lookup all (klass, signed claims: SignedClaims) -> Query | None:
    # This lookup requires the environment claim to be present;
    # if it isn't, bail out early.
    if not (environment := signed_claims.get("environment")):
        return None
    repository = signed_claims["repository"]
    repository_owner, repository_name = repository.split("/", 1)
    workflow_prefix = f"{repository}/.github/workflows/"
    workflow_ref = signed_claims["job_workflow_ref"].removeprefix(workflow_prefix)
    return (
        Query(klass)
        .filter_by(
            repository_name=repository_name,
            repository owner=repository owner,
            repository_owner_id=signed_claims["repository_owner_id"],
            environment=environment.lower(),
        .filter(
            literal(workflow ref).like(func.concat(klass.workflow filename, "%"))
```

```
@staticmethod
def __lookup_no_environment__(klass, signed_claims: SignedClaims) -> Query | None:
   repository = signed_claims["repository"]
   repository owner, repository name = repository.split("/", 1)
    workflow_prefix = f"{repository}/.github/workflows/"
   workflow ref = signed claims["job workflow ref"].removeprefix(workflow prefix)
   return (
        Query(klass)
        .filter by(
            repository_name=repository_name,
            repository_owner=repository_owner,
           repository_owner_id=signed_claims["repository_owner_id"],
           environment=None,
         filter(
             iteral(workflow ref).like(func.concat(klass.workflow filename, "%"))
```

terrible!

```
__lookup_strategies__ = [
    __lookup_all__,
    __lookup_no_environment__,
```



## Lesson #3: OIDC tokens are not API tokens

"An OIDC credential is basically an API token; I don't need to do an exchange"

Reality: you can do this, but...

- OIDC credentials aren't plugged into your preexisting AuthN/Z or permissions/scopes; you'll end up re-implementing a bunch of what you already have (and reimplementation means more bugs)
- OIDC credentials are chonky contain all kinds of stuff you might not want to hold onto for prolonged periods (user emails, other potential PII)
- IdPs can change expiration and other policies without notice; creating your own temporary token makes you resilient to these changes!

Conclusion: Performing token exchange minimizes the amount of novel code needed; reduces potential sources of PII; offers additional resilience against IdP changes.

### The name is inaccurate and even misleading i

### Lesson #4: Words are hard

This is a great feature, but its name is going to confuse people.

Definitely a step in the right direction, but not what I thought it was going to be given the name

This is a pretty misleading title.

The name of this program is misleading.

Have to say this is a REALLY misleading name. '

### Lesson #5: It's all worth it!

OIDC is complicated; trusted publishing's model on top of it even more so.

But the user experience and security gains make it worth it:

- User feedback (once they understand it) is overwhelmingly positive
  - High demand for more IdPs/more trust relationships (e.g. reusable GitHub workflows)
- Adoption by critical projects is (slowly) advancing:
  - o 219 critical projects have trusted publishers, 152 actively publishing (2023-09-18)
  - High and growing "top 20" coverage: urllib3, certifi, cryptography, charset-normalizer, idna, wheel, etc.



### **Alex Gaynor**

I'm one of the primary authors of the one of the first packages to adopt the OIDC support in PyPI (package includes Rust code, naturally). Happy to gush about how great it was :-)



#### The future

### What comes next?

The same building blocks that give us trusted publishing (OIDC, machine identities) are also the building blocks for build provenance and code signing!

The goal: ecosystems that support trusted publishing should find it relatively easy to enable Sigstore for codesigning.

#### **Event**

#### File added to version 2.6.1

Filename: pip\_audit-2.6.1.tar.gz

Added by: OpenID created token URL: https://github.com/pypa/pip-

audit/commit/d4242095300357730e1510ef2db837cf1b1142f1

### File added to version 2.6.1

Filename: pip\_audit-2.6.1-py3-none-any.whl

Added by: OpenID created token URL: https://github.com/pypa/pip-

audit/commit/d4242095300357730e1510ef2db837cf1b1142f1

#### Version 2.6.1 created

Added by: OpenID created token

URL: https://github.com/pypa/pip-

audit/commit/d4242095300357730e1510ef2db837cf1b1142f1



#### Conclusion

# Takeaways

- Trusted publishing is a double win: both for security and for usability
  - The best kind of security improvements make users' lives easier, not harder!
  - Cynically: the *only* kinds of security improvements that matter are the ones that engineers want to use
- Trusted publishing is not tied to PyPI; other package indices can use the same techniques and reap the same benefits!
  - We (Trail of Bits) would be thrilled to reapply our experience on PyPI to other ecosystems; please come find me during the day and chat with me about it!
- Trusted publishing is a logical step towards our shared supply chain goals: source and build provenance, code signing, generalized verifiable attestations over software/dependency graphs

#### end of the talk

# thank you!

these slides will soon be available here:

### https://yossarian.net/publications#ossfeu-2023

#### resources:

- docs.pypi.org/trusted-publishers: official PyPI documentation for trusted publishing
- "Introducing 'Trusted Publishers": official PyPI announcement post
- "Trusted publishing: a new benchmark for packaging security": ToB's writeup on trusted publishing, threat modeling, etc.

### **Contact:**

- william@trailofbits.com
- @yossarian@infosec.exchange 0